

✓ II. Implement the following:

(Implementation 5 marks and Visualization and documentation 5 marks) • Scenario: The XOR gate is known for its complexity, as it outputs 1 only when the inputs are different. This is a challenge for a Single Layer Perceptron since XOR is not linearly separable. • Lab Task: Attempt to implement a Single Layer Perceptron in Google Colab to classify the output of an XOR gate. Perform the following steps: • Create the XOR gate's truth table dataset. • Implement the perceptron model and train it using the XOR dataset using MCP (McCulloch Pitts) Neuron. • Observe and discuss the perceptron's performance in this scenario. • Implement XOR using Multi-Layer Perceptron.

```
import numpy as np

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([0, 1, 1, 0]) # target values

weights = np.random.rand(2)
bias = np.random.rand(1)
learning_rate = 0.1

def step_function(x):
    return np.where(x >= 0, 1, 0)

def perceptron_train(X, y, weights, bias, learning_rate, epochs=10):
    for epoch in range(epochs):
        for i in range(len(X)):
            net_input = np.dot(X[i], weights) + bias
            output = step_function(net_input)
            error = y[i] - output
            weights += learning_rate * error * X[i]
            bias += learning_rate * error
    return weights, bias

weights, bias = perceptron_train(X, y, weights, bias, learning_rate)

def perceptron_predict(X, weights, bias):
    net_input = np.dot(X, weights) + bias
    return step_function(net_input)

predictions = perceptron_predict(X, weights, bias)
print("Predictions after training:", predictions)
```

↗ Predictions after training: [1 1 0 0]

✓ Discuss the perceptron's performance

The perceptron is not be able to correctly classify all XOR inputs because XOR is not linearly separable.

```
from sklearn.neural_network import MLPClassifier

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([0, 1, 1, 0])

mlp = MLPClassifier(hidden_layer_sizes=(2,), activation='relu', solver='adam', max_iter=10000)
mlp.fit(X, y)

predictions_mlp = mlp.predict(X)
print("MLP Predictions:", predictions_mlp)
```

↗ MLP Predictions: [0 1 1 0]

Multi-Layer Perceptron adding a hidden layer, the MLP is capable of classifying XOR correctly,because it can understand the hidden layer ,which help to learn all the layers and get the output correctly.

✓ A. Sentiment Analysis Twitter Airline

Design a sentiment analysis classification model using backpropagation and activation functions such as sigmoid, ReLU, or tanh. Implement a neural network that can classify sentiment (positive/negative) from a small dataset. Demonstrate how backpropagation updates the weights during the training process. (link Provided at the top of the page to download the dataset) Task: • Create a simple feed-forward neural network for binary sentiment classification (positive/negative). • Use backpropagation to optimize the model's weights based on error calculation. • Experiment with different activation functions (sigmoid, ReLU, tanh) in the hidden layer and compare the model's performance. • Evaluate the model on a test set using accuracy and plot the loss over epochs.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer

dataset = '/content/Tweets - Tweets.csv'
tweets_df = pd.read_csv(dataset)

filtered_df = tweets_df[['airline_sentiment', 'text']].copy()

filtered_df = filtered_df[filtered_df['airline_sentiment'].isin(['positive', 'negative'])]

filtered_df['label'] = filtered_df['airline_sentiment'].map({'positive': 1, 'negative': 0})

filtered_df = filtered_df[['text', 'label']].dropna()

vectorizer = TfidfVectorizer(stop_words='english', max_features=1000)

X = vectorizer.fit_transform(filtered_df['text'])

y = filtered_df['label'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Training Data Shape:", X_train.shape)
print("Testing Data Shape:", X_test.shape)
```

→ Training Data Shape: (9232, 1000)
Testing Data Shape: (2309, 1000)

Activation functions

```
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def relu(x):
    return np.maximum(0, x)

def tanh(x):
    return np.tanh(x)

def sigmoid_derivative(x):
    return sigmoid(x) * (1 - sigmoid(x))

def relu_derivative(x):
    return np.where(x > 0, 1, 0)

def tanh_derivative(x):
    return 1 - np.tanh(x) ** 2

def initialize_weights(input_size, hidden_size, output_size):
    W1 = np.random.randn(input_size, hidden_size)
    b1 = np.zeros((1, hidden_size))
    W2 = np.random.randn(hidden_size, output_size)
    b2 = np.zeros((1, output_size))
    return W1, b1, W2, b2

def forward_propagation(X, W1, b1, W2, b2, activation_func):
    Z1 = np.dot(X, W1) + b1
    A1 = activation_func(Z1)
    Z2 = np.dot(A1, W2) + b2
    A2 = sigmoid(Z2)
    return Z1, A1, Z2, A2

def backpropagation(X, y, Z1, A1, A2, W2, activation_derivative):
    m = X.shape[0]
```

```

dZ2 = A2 - y
dW2 = np.dot(A1.T, dZ2) / m
db2 = np.sum(dZ2, axis=0, keepdims=True) / m

dA1 = np.dot(dZ2, W2.T)
dZ1 = dA1 * activation_derivative(Z1)
dW1 = np.dot(X.T, dZ1) / m
db1 = np.sum(dZ1, axis=0, keepdims=True) / m

return dW1, db1, dW2, db2

def update_weights(W1, b1, W2, b2, dW1, db1, dW2, db2, learning_rate):
    W1 -= learning_rate * dW1
    b1 -= learning_rate * db1
    W2 -= learning_rate * dW2
    b2 -= learning_rate * db2
    return W1, b1, W2, b2

def train_neural_network(X_train, y_train, hidden_size, epochs, learning_rate, activation_func, activation_derivative):
    input_size = X_train.shape[1]
    output_size = 1
    W1, b1, W2, b2 = initialize_weights(input_size, hidden_size, output_size)

    for epoch in range(epochs):

        Z1, A1, Z2, A2 = forward_propagation(X_train, W1, b1, W2, b2, activation_func)

        loss = -np.mean(y_train * np.log(A2) + (1 - y_train) * np.log(1 - A2))

        dW1, db1, dW2, db2 = backpropagation(X_train, y_train, Z1, A1, A2, W2, activation_derivative)

        W1, b1, W2, b2 = update_weights(W1, b1, W2, b2, dW1, db1, dW2, db2, learning_rate)

        if epoch % 100 == 0:
            print(f'Epoch {epoch}, Loss: {loss:.4f}')

    return W1, b1, W2, b2

def predict(X, W1, b1, W2, b2, activation_func):
    _, _, A2 = forward_propagation(X, W1, b1, W2, b2, activation_func)
    return (A2 > 0.5).astype(int)

hidden_size = 64
epochs = 1000
learning_rate = 0.01

activation_func = relu
activation_derivative = relu_derivative

W1, b1, W2, b2 = train_neural_network(X_train.toarray(), y_train.reshape(-1, 1), hidden_size, epochs, learning_rate, activation_func, activation_derivative)

predictions = predict(X_test.toarray(), W1, b1, W2, b2, activation_func)

accuracy = np.mean(predictions == y_test.reshape(-1, 1))
print(f"Test Accuracy: {accuracy:.4f}")

→ Epoch 0, Loss: 1.8556
Epoch 100, Loss: 1.3104
Epoch 200, Loss: 1.2197
Epoch 300, Loss: 1.1505
Epoch 400, Loss: 1.0892
Epoch 500, Loss: 1.0342
Epoch 600, Loss: 0.9853
Epoch 700, Loss: 0.9412
Epoch 800, Loss: 0.9018
Epoch 900, Loss: 0.8665
Test Accuracy: 0.7181

```

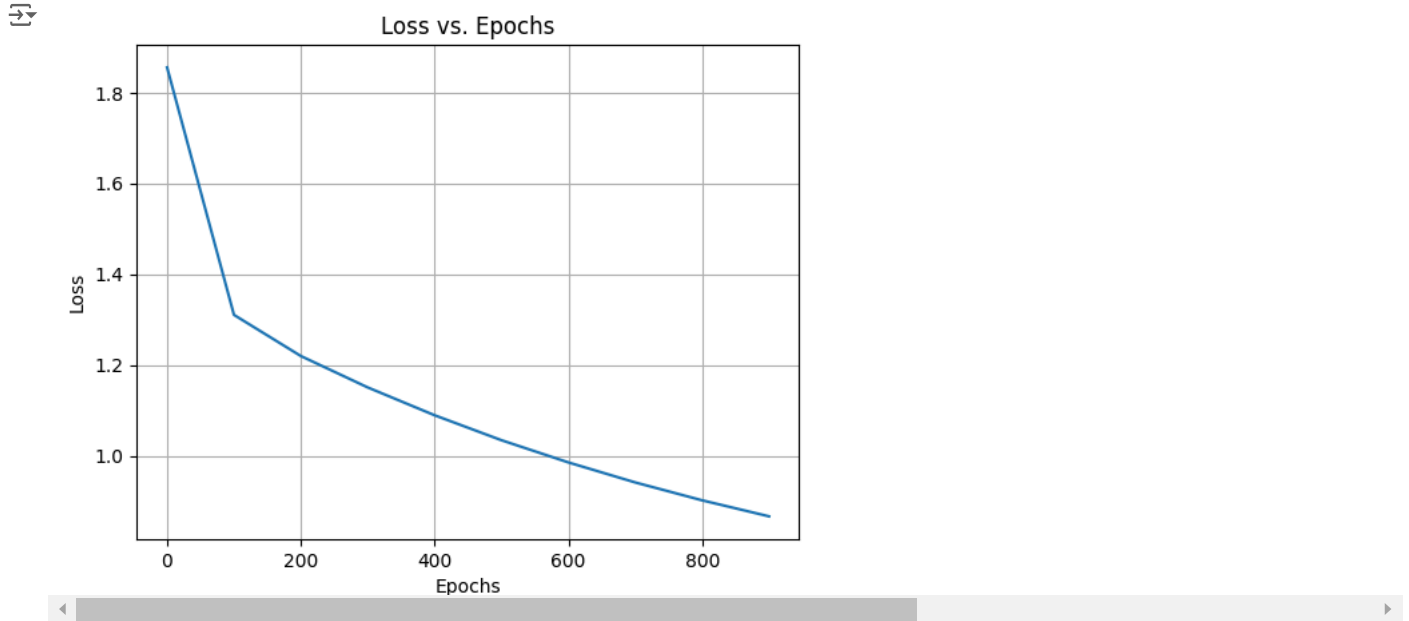
Backproagation

```

import matplotlib.pyplot as plt
epochs = [0, 100, 200, 300, 400, 500, 600, 700, 800, 900]
losses = [1.8556, 1.3104, 1.2197, 1.1505, 1.0892, 1.0342, 0.9853, 0.9412, 0.9018, 0.8665]

```

```
plt.plot(epochs, losses)
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.grid(True)
plt.show()
```



```
hidden_size = 64
epochs = 1000
learning_rate = 0.01
```

```
activation_func = sigmoid
activation_derivative = sigmoid_derivative
```

```
W1, b1, W2, b2 = train_neural_network(X_train.toarray(), y_train.reshape(-1, 1), hidden_size, epochs, learning_rate, activation_func, ac
```

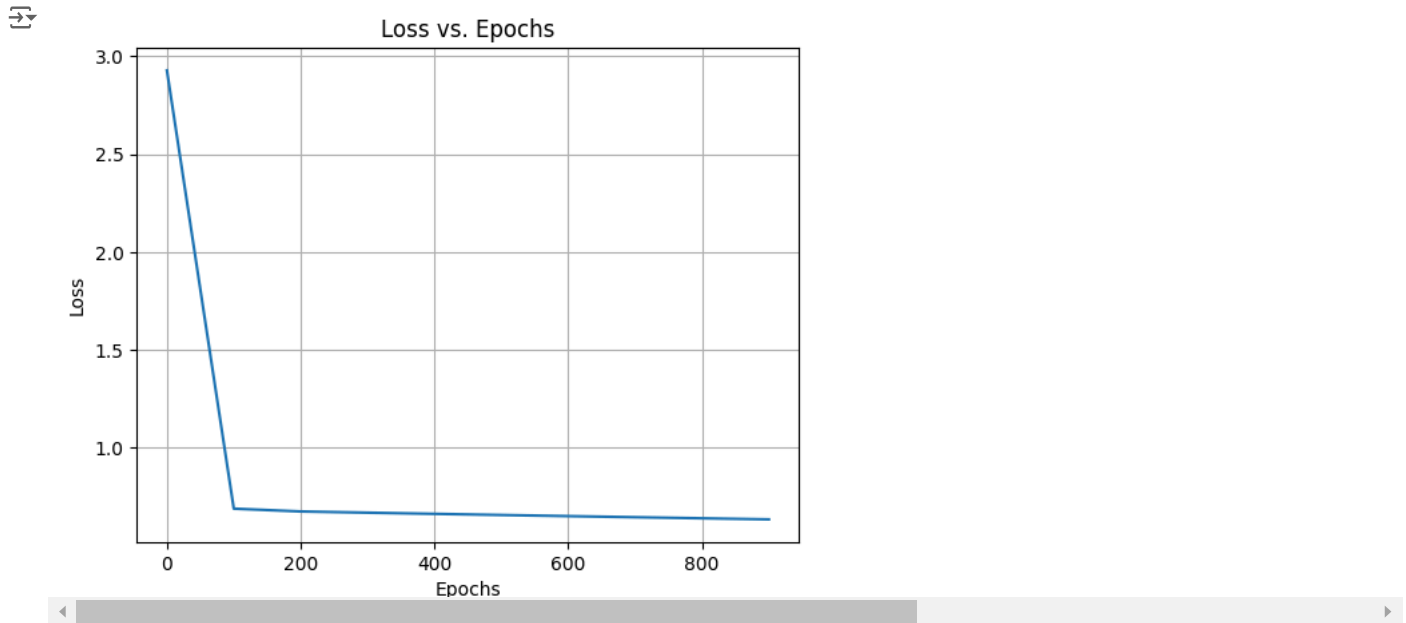
```
predictions = predict(X_test.toarray(), W1, b1, W2, b2, activation_func)
```

```
accuracy = np.mean(predictions == y_test.reshape(-1, 1))
print(f"Test Accuracy: {accuracy:.4f}")
```

```
Epoch 0, Loss: 2.9272
Epoch 100, Loss: 0.6902
Epoch 200, Loss: 0.6759
Epoch 300, Loss: 0.6697
Epoch 400, Loss: 0.6637
Epoch 500, Loss: 0.6579
Epoch 600, Loss: 0.6521
Epoch 700, Loss: 0.6465
Epoch 800, Loss: 0.6410
Epoch 900, Loss: 0.6356
Test Accuracy: 0.7631
```

```
epochs = [0, 100, 200, 300, 400, 500, 600, 700, 800, 900]
losses = [2.9272, 0.6902, 0.6759, 0.6697, 0.6637, 0.6579, 0.6521, 0.6465, 0.6410, 0.6356]
```

```
plt.plot(epochs, losses)
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.grid(True)
plt.show()
```



using relu activation function the accuracy is better sigmoid because it helps in vanishing gradient problem. where as in sigmoid it tends to vanish the gradient.